



Block 3

Introduction to Relational Database Management Systems

Learning Outcomes

After completing this topic and the recommended reading, you should be able to:

- Use relational database models and structured query languages (SQL).
- Gain experience with interfacing SQL from R and Python.

1. Database

Database

- An organised collection of structured information or data.
- Stored electronically in a computer system.

Relational Databases

- A type of database that stores and provides access to data points that are related to one another.
- Store data in structured, tabular form.
- The columns of the table hold **attributes** of the data.
- The rows of the table hold **records (tuples)**, and each record usually has a value for each attribute.

Nonrelational Databases

- Often called **NoSQL** databases.
- Store data in unstructured and semi-structured, non-tabular form.
- Graph database
 - Stores data in terms entities (nodes) and relationships (edges) between entities.
- Document-oriented database
 - Stores data in the form of JSON-like documents.
 - Stores all information for a given object in a single instance in the database.
 - Every stored object can be different from another.
- Examples:
 - *MongoDB*: uses JSON-like documents with optional schemas.

Database Management System (DBMS)

- A database is usually controlled by a DBMS.
- The data can then be easily accessed, managed, modified, updated, controlled and organised.

Relational Database Management System (RDBMS)

- A system used to maintain relational databases.
- Stores data in a row-based table structure which connects related data elements.
- Using **Structured Query Language (SQL)** to access the database.
- Includes functions that maintain the security, accuracy, integrity and consistency of the data.

RDBMS Terminology

- **Relation (Table)**
 - Collection of rows and columns.
 - Each table usually represent an entity.
- **Attribute (Column)**
 - Each attribute has a type or domain.
- **Tuple (Row)**
 - Each tuple represent a record, a set of attribute values.
- **Schema**
 - The description on how the database table is constructed.

- **Primary Key**

- The attribute which is the unique identifier (**ID**) for each tuple/record.
- The set of attributes whose combine values are unique.
- Cannot have null values.

- **Foreign Key**

- The primary key that is used in another table that provides a link/relationship between data in two tables.

Table: **Course**

	course_id	name	capacity
	Filter	Filter	Filter
1	ST101	programming for data science	60
2	ST115	Managing and Visualising Data	60
3	ST207	Databases	30
4	ST310	Machine Learning	100

Table: **Student**

	student_id	name	year
	Filter	Filter	Filter
1	201921323	Ava Smith	2
2	201832220	Ben Johnson	3
3	202003219	Charlie Jones	1
4	202045234	Dan Norris	1
5	201985603	Emily Wood	1
6	201933222	Freddie Harris	2
7	201875940	Grace Clarke	2

Table: **Grade**

	course_id	student_id	final_mark
	Filter	Filter	Filter
1	ST101	201921323	78
2	ST101	201985603	60
3	ST101	202003219	47
4	ST115	201921323	92
5	ST115	202003219	67
6	ST115	201933222	88
7	ST207	201933222	73
8	ST207	201875940	60

Name	Type	Schema
Tables (3)		
Course		CREATE TABLE "Course" ("course_id" TEXT, "name" TEXT, "capacity" INTEGER)
course_id	TEXT	"course_id" TEXT
name	TEXT	"name" TEXT
capacity	INTEGER	"capacity" INTEGER
Grade		CREATE TABLE "Grade" ("course_id" TEXT, "student_id" INTEGER, "final_mark" INTEGER)
course_id	TEXT	"course_id" TEXT
student_id	INTEGER	"student_id" INTEGER
final_mark	INTEGER	"final_mark" INTEGER
Student		CREATE TABLE "Student" ("student_id" TEXT, "name" TEXT, "year" INTEGER, PRIMARY KEY("student_id"))
student_id	TEXT	"student_id" TEXT
name	TEXT	"name" TEXT
year	INTEGER	"year" INTEGER

2. Structured Query Language

Structured Query Language (SQL)

- A standardised domain-specific language used in programming and designed for managing data held in a RDBMS.
- Used to create, maintain (insert, update, delete), and retrieve (query) the relational database.

SQLite

- RDBMS contained in a C library.
- Lightweight, non-client-server database engine.

DB Browser for SQLite (DB4S)

- Tool to create, design, and edit database files that are compatible with SQLite.
- <https://sqlitebrowser.org>

Basic SQLite Syntax (Creating & Manipulating Databases)

- **Add a Table**
 - `CREATE TABLE table_name (`
`attribute1 datatype,`
`attribute2 datatype,`
`...`

```
1  CREATE TABLE Teacher (
2    staff_id TEXT,
3    name TEXT
4 )
```

- **Delete a Table**

- `DROP TABLE table_name`

```
1  DROP TABLE Teacher
```

- **Select Data**

- `SELECT attribute1, attribute2, ...
FROM table_name`

```
1  SELECT name  
2  FROM Student
```

- **Insert Tuples/Rows**

- `INSERT INTO table_name (attribute1, attribute2, ...)
VALUES (value1, value2, ...)`
 - `INSERT INTO table_name
VALUES (value1, value2, ...)`

```
1  INSERT INTO Student  
2  VALUES (202029744, "Harper Taylor", 1)
```

- **Update Tuples/Rows**

- `UPDATE table_name
SET attribute1 = value1, attribute2 = value2, ...
WHERE conditions`

```
1  UPDATE Student  
2  SET student_id = "201929744"  
3  WHERE name = "Harper Taylor"
```

- **Delete Tuples/Rows**

- `DELETE FROM table_name
WHERE conditions`

```

1  DELETE FROM Student
2  WHERE name = "Harper Taylor"

```

Basic SQLite Queries

- Conditions

- SELECT *attribute1, attribute2, ...*
 FROM *table_name*
 WHERE *conditions*

```

1  SELECT student_id
2  FROM Grade
3  WHERE course_id = 'ST101'

```

```

1  SELECT *
2  FROM Grade
3  WHERE course_id = 'ST101'

```

- Several Tables

- SELECT *table.attribute1, table.attribute2, ...*
 FROM *table1, table2, ...*
 WHERE *condition1 AND/OR condition2 ...*

```

1  SELECT Student.name
2  FROM Grade, Student
3  WHERE Grade.course_id = 'ST101' AND Student.student_id = Grade.student_id

```

```

1  SELECT name
2  FROM Grade, Student
3  WHERE Grade.course_id = 'ST101' AND Student.student_id = Grade.student_id

```

- Multiple Conditions

```

1  SELECT Course.name
2  FROM Student, Grade, Course
3  WHERE (Student.name = 'Ava Smith' OR Student.name = 'Freddie Harris')
4          AND Student.student_id = Grade.student_id
5          AND Course.course_id = Grade.course_id

```

- SELECT DISTINCT *attribute1*, *attribute2*, ...

FROM *table_name*

```
1  SELECT DISTINCT Course.name
2  FROM Student, Grade, Course
3  WHERE (Student.name = 'Ava Smith' OR Student.name = 'Freddie Harris')
4      AND Student.student_id = Grade.student_id
5      AND Course.course_id = Grade.course_id
```

- Aggregation

- SELECT *attributes*, *aggregation_functions* AS *column_name*

FROM *table_name*

WHERE *conditions*

GROUP BY *attributes*

ORDER BY *attributes*

```
1  SELECT course_id, AVG(final_mark)
2  FROM Grade
3  GROUP BY course_id
```

```
1  SELECT course_id, AVG(final_mark) AS avg_mark
2  FROM Grade
3  GROUP BY course_id
```

- Aggregation functions

- COUNT(), MAX(), MIN(), SUM(), AVG()

Basic SQLite Joins

- A ***JOIN*** clause is used to combine rows from two or more tables, based on a related column between them.
- **Inner Join**
 - Selects records that have matching values in both tables

```
1  SELECT *
2  FROM Grade, Student
3  WHERE Grade.course_id = 'ST101' AND Student.student_id = Grade.student_id
4  ORDER BY Student.name
```

- SELECT *attributes*

FROM *table1* JOIN *table2*
ON *table1.attribute* = *table2.attribute*

```
1  SELECT *
2  FROM Student JOIN Grade
3  ON Student.student_id = Grade.student_id
4  WHERE course_id = 'ST101'
5  ORDER BY Student.name
```

- SELECT *attributes*

FROM *table1* JOIN *table2*
USING(*attribute*)

```
1  SELECT *
2  FROM Student JOIN Grade
3  USING(student_id)
4  WHERE course_id = 'ST101'
5  ORDER BY Student.name
```

- Natural Join

- The join condition is automatically identified
- SELECT *attributes*

FROM *table1* NATURAL JOIN *table2*
WHERE *conditions*
ORDER BY *attribute1*

```
1  SELECT *
2  FROM Student NATURAL JOIN Grade
3  WHERE course_id = 'ST101'
4  ORDER BY Student.name
```

- **Left Join**

- Returns all records from the left table (table1), and the matching records from the right table (table2)
- *NULL* value for right table (table2) attributes with no corresponding record
- *SELECT attributes*

FROM *table1* LEFT JOIN *table2*
USING (*attribute*)

```
1  SELECT *
2  FROM Student LEFT JOIN Grade
3  USING (student_id)
4  ORDER BY Student.name
```

- **Cross Join**

- Returns all records when there is a match in left table (table1) or right table (table2) records
- *SELECT attributes*

FROM *table1* CROSS JOIN *table2*
WHERE *conditions*
ORDER BY *attribute1*

```
1  SELECT *
2  FROM Student CROSS JOIN Grade
3  ORDER BY Student.name
```

3. Using Databases with R



Connecting to Databases in R (DBI)

- Use “RSQLite” and “DBI” packages
 - `install.packages("RSQLite")`
 - `library(DBI)`
- Set working directory
 - `setwd("~/ST2195/database")`
- Remove existing database
 - `if (file.exists("University.db"))`
`file.remove("University.db")`
- Create connection to database
 - `conn <- dbConnect(RSQLite::SQLite(), "University.db")`

Creating Tables in R (DBI)

- Read CSV files as data frames
 - `course <- read.csv("course.csv", header = TRUE)`

```
course_id, name, capacity
"ST101", "programming for data science", 60
"ST115", "Managing and Visualising Data", 60
"ST207", "Databases", 30
"ST310", "Machine Learning", 100
```

- `student <- read.csv("student.csv", header = TRUE)`

student_id	name	year
201921323	Ava Smith	2
201832220	Ben Johnson	3
202003219	Charlie Jones	1
202045234	Dan Norris	1
201985603	Emily Wood	1
201933222	Freddie Harris	2
201875940	Grace Clarke	2

- o grade <- read.csv("grade.csv", header = TRUE)

course_id	student_id	final_mark
ST101	201921323	78
ST101	201985603	60
ST101	202003219	47
ST115	201921323	92
ST115	202003219	67
ST115	201933222	88
ST207	201933222	73
ST207	201875940	60

- Copy data frames to database tables
 - o dbWriteTable(conn, "Course", course)
 - o dbWriteTable(conn, "Student", student)
 - o dbWriteTable(conn, "Grade", grade)
- Adding a new table
 - o dbCreateTable(conn, "Teacher",


```
c(staff_id = "TEXT", name = "TEXT"))
```

 or
 - o dbExecute(conn, "CREATE TABLE Teacher (
 staff_id TEXT PRIMARY KEY,
 name TEXT)")
- Deleting a table
 - o dbRemoveTable(conn, "Teacher")

or

- dbExecute(conn, “DROP TABLE Teacher”)
- List database tables
 - dbListTables(conn)

```
[1] "Course"  "Grade"   "Student"
```

- Browse database table
 - dbReadTable(conn, “Student”)

	student_id	name	year
1	201921323	Ava Smith	2
2	201832220	Ben Johnson	3
3	202003219	Charlie Jones	1
4	202045234	Dan Norris	1
5	201985603	Emily Wood	1
6	201933222	Freddie Harris	2
7	201875940	Grace Clarke	2

- Show attributes
 - dbListFields(conn, “Student”)

```
[1] "student_id" "name"      "year"
```

Manipulating Data in R (DBI)

- Inserting tuples/rows
 - dbAppendTable(conn, “Student”,
 data.frame(student_id = “202029744”,
 name = “Harper Taylor”, year = 1))

or

- dbExecute(conn, “INSERT INTO Student
VALUES(202029744, ‘Harper Taylor’, 1)”)

	student_id	name	year
1	201921323	Ava Smith	2
2	201832220	Ben Johnson	3
3	202003219	Charlie Jones	1
4	202045234	Dan Norris	1
5	201985603	Emily Wood	1
6	201933222	Freddie Harris	2
7	201875940	Grace Clarke	2
8	202029744	Harper Taylor	1

- Updating tuples/rows

- dbExecute(conn, “UPDATE Student
SET student_id = ‘201929744’
WHERE name = ‘Harper Taylor’”)

	student_id	name	year
1	201921323	Ava Smith	2
2	201832220	Ben Johnson	3
3	202003219	Charlie Jones	1
4	202045234	Dan Norris	1
5	201985603	Emily Wood	1
6	201933222	Freddie Harris	2
7	201875940	Grace Clarke	2
8	201929744	Harper Taylor	1

- Deleting tuples/rows

- dbExecute(conn, “DELETE FROM Student
WHERE name = ‘Harper Taylor’”)

	student_id	name	year
1	201921323	Ava Smith	2
2	201832220	Ben Johnson	3
3	202003219	Charlie Jones	1
4	202045234	Dan Norris	1
5	201985603	Emily Wood	1
6	201933222	Freddie Harris	2
7	201875940	Grace Clarke	2

- Disconnecting from database
 - dbDisconnect(conn)

Querying Databases in R (DBI)

- Getting query result in data frame
 - ```
q1 <- dbGetQuery(conn, "SELECT final_mark
 FROM Grade
 WHERE course_id = 'ST101'")
```

|   | final_mark |
|---|------------|
| 1 | 78         |
| 2 | 60         |
| 3 | 47         |

- Sending query to database engine
  - ```
q1 <- dbSendQuery(conn, "SELECT final_mark
                           FROM Grade
                           WHERE course_id = 'ST101'")
```

```
<SQLiteResult>
  SQL  SELECT final_mark
  FROM Grade
  WHERE course_id = 'ST101'
  ROWS Fetched: 0 [incomplete]
  Changed: 0
```

- Fetch query result from database engine
 - dbFetch(q1)

	final_mark
1	78
2	60
3	47

- Getting results in alphabetical order
 - dbGetQuery(conn, “SELECT Student.name
FROM Grade, Student
WHERE Grade.course_id = ‘ST101’ AND
Student.student_id = Grade.student_id
ORDER BY Student.name”)
 - or
 - dbGetQuery(conn, “SELECT Student.name
FROM Student NATURAL JOIN Grade
WHERE course_id = ‘ST101’
ORDER BY Student.name”)

	name
1	Ava Smith
2	Charlie Jones
3	Emily Wood

- Getting distinct results

- dbGetQuery(conn, “SELECT DISTINCT Course.name
FROM Student, Grade, Course
WHERE (Student.name = ‘Ava Smith’ OR
Student.name = ‘Freddie Harris’) AND
Student.student_id = Grade.student_id
AND Course.course_id = Grade.course_id”)
or
○ dbGetQuery(conn, “SELECT DISTINCT Course.name
FROM (Student NATURAL JOIN Grade)
S JOIN Course
ON Course.course_id = S.course_id
WHERE S.name = ‘Ava Smith’ OR
S.name = ‘Freddie Harris’”)

	name
1	programming for data science
2	Managing and Visualising Data
3	Databases

- Getting calculated results

- dbGetQuery(conn, “SELECT course_id, AVG(final_mark)
AS avg_mark
FROM Grade
GROUP BY course_id”)

	course_id	avg_mark
1	ST101	61.66667
2	ST115	82.33333
3	ST207	66.50000

Querying Databases in R (dplyr)

- Creating a reference to tables
 - `library(dplyr)`
 - `student_db <- tbl(conn, "Student")`
 - `grade_db <- tbl(conn, "Grade")`
 - `course_db <- tbl(conn, "Course")`
- Getting rows with conditions
 - `q1 <- grade_db %>% filter(course_id == "ST101")`

course_id	student_id	final_mark
<chr>	<int>	<int>
1 ST101	201921323	78
2 ST101	201985603	60
3 ST101	202003219	47
 - `show_query(q1)`

```

<SQL>
SELECT *
FROM `Grade`
WHERE (`course_id` = 'ST101')

```
- Getting rows in alphabetical order
 - `q2 <- inner_join(student_db, grade_db) %>%
 filter(course_id == "ST101") %>%
 select(name) %>%
 arrange(name)`

```

# Ordered by: name
name
<chr>
1 Ava Smith
2 Charlie Jones
3 Emily Wood

```

- o show_query(q2)

```
<SQL>
SELECT `name`
FROM (SELECT `LHS`.`student_id` AS `student_id`, `name`, `year`, `course_id`, `final_mark`
      FROM `Student` AS `LHS`
      INNER JOIN `Grade` AS `RHS`
      ON (`LHS`.`student_id` = `RHS`.`student_id`)
     )
WHERE (`course_id` = 'ST101')
ORDER BY `name`
```

- Getting distinct rows

- o q3 <- inner_join(student_db, grade_db, by = “student_id”) %>%
 inner_join(course_db, by = “course_id”,
 suffix = c(“.student”, “.course”) %>%
 filter(name.student == ‘Ava Smith’ |
 name.student == ‘Freddie Harris’) %>%
 select(name.course) %>%
 distinct()

```
name.course
<chr>
1 programming for data science
2 Managing and Visualising Data
3 Databases
```

- o show_query(q3)

```
<SQL>
SELECT DISTINCT `name.course`
FROM (SELECT `student_id`, `LHS`.`name` AS `name.student`, `year`, `LHS`.`course_id` AS `course_id`, `final_mark`, `RHS`.`name` AS `name.course`, `capacity`
      FROM (SELECT `LHS`.`student_id` AS `student_id`, `name`, `year`, `course_id`, `final_mark`
            FROM `Student` AS `LHS`
            INNER JOIN `Grade` AS `RHS`
            ON (`LHS`.`student_id` = `RHS`.`student_id`)
           ) AS `LHS`
      INNER JOIN `Course` AS `RHS`
      ON (`LHS`.`course_id` = `RHS`.`course_id`)
     )
WHERE (`name.student` = 'Ava Smith' OR `name.student` = 'Freddie Harris')
```

- Getting calculated rows

- ```
q4 <- grade_db %>% group_by(course_id) %>%
 summarize(avg_mark = mean(final_mark, na.rm = TRUE))
```

```
course_id avg_mark
<chr> <dbl>
1 ST101 61.7
2 ST115 82.3
3 ST207 66.5
```

- `show_query(q4)`

```
<SQL>
SELECT `course_id`, AVG(`final_mark`) AS `avg_mark`
FROM `Grade`
GROUP BY `course_id`
```

## 4. Using Databases with Python



### *Connecting to Databases in Python*

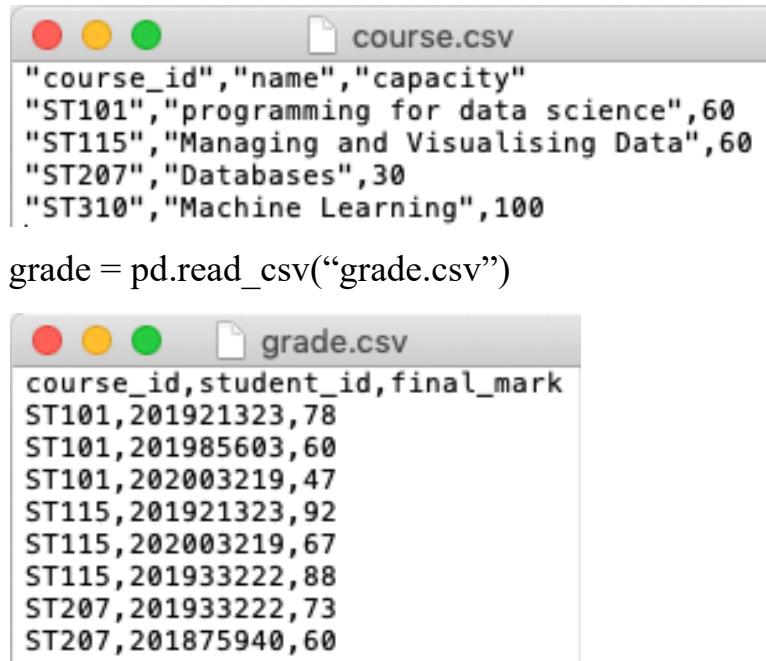
- Remove existing database
  - import os
  - try:
  - os.remove('University.db')
  - except OSError:
  - pass
- Use “sqlite3” packages
  - import sqlite3
- Create connection to database
  - conn = sqlite3.connect('University.db')

### *Creating Tables in Python*

- Read CSV files as data frames
  - import pandas as pd
  - student = pd.read\_csv("student.csv")

student.csv	
student_id	name,year
201921323	Ava Smith,2
201832220	Ben Johnson,3
202003219	Charlie Jones,1
202045234	Dan Norris, 1
201985603	Emily Wood,1
201933222	Freddie Harris,2
201875940	Grace Clarke,2

- course = pd.read\_csv("course.csv")



The image shows a Jupyter Notebook interface with two code cells. The first cell displays the content of 'course.csv' with the following data:

```
"course_id", "name", "capacity"
"ST101", "programming for data science", 60
"ST115", "Managing and Visualising Data", 60
"ST207", "Databases", 30
"ST310", "Machine Learning", 100
```

The second cell displays the content of 'grade.csv' with the following data:

```
course_id, student_id, final_mark
ST101, 201921323, 78
ST101, 201985603, 60
ST101, 202003219, 47
ST115, 201921323, 92
ST115, 202003219, 67
ST115, 201933222, 88
ST207, 201933222, 73
ST207, 201875940, 60
```

- Copy data frames to database tables
  - `student.to_sql('Student', con = conn, index = False)`
  - `course.to_sql('Course', con = conn, index = False)`
  - `grade.to_sql('Grade', con = conn, index = False)`

## ***Manipulating Data in Python***

- Create cursor object
  - `c = conn.cursor()`
- Execute SQL commands to get all tables
  - `c.execute('' SELECT name  
FROM sqlite_master  
WHERE type='table'  
'')`
- Fetch all tables
  - `c.fetchall()`

```
[('Student',), ('Course',), ('Grade',)]
```

- Browse database table

- `q = c.execute("SELECT * FROM Student").fetchall()`
- `pd.DataFrame(q)`

	0	1	2
0	201921323	Ava Smith	2
1	201832220	Ben Johnson	3
2	202003219	Charlie Jones	1
3	202045234	Dan Norris	1
4	201985603	Emily Wood	1
5	201933222	Freddie Harris	2
6	201875940	Grace Clarke	2

- Add a new table

- `c.execute(`` CREATE TABLE Teacher  
(staff_id TEXT PRIMARY KEY, name TEXT)  
```)`
- `conn.commit()`

- Delete a table

- `c.execute("DROP TABLE Teacher")`
- `conn.commit()`

- Insert tuples/rows

- `c.execute(`` INSERT INTO Student
VALUES(202029744, 'Harper Taylor', 1)
```)`
- `conn.commit()`

	0	1	2
0	201921323	Ava Smith	2
1	201832220	Ben Johnson	3
2	202003219	Charlie Jones	1
3	202045234	Dan Norris	1
4	201985603	Emily Wood	1
5	201933222	Freddie Harris	2
6	201875940	Grace Clarke	2
7	202029744	Harper Taylor	1

- Update tuples/rows

- `c.execute("“ UPDATE Student  
SET student_id = “201929744”  
WHERE name = “Harper Taylor”  
”“)`

	0	1	2
0	201921323	Ava Smith	2
1	201832220	Ben Johnson	3
2	202003219	Charlie Jones	1
3	202045234	Dan Norris	1
4	201985603	Emily Wood	1
5	201933222	Freddie Harris	2
6	201875940	Grace Clarke	2
7	201929744	Harper Taylor	1

- Delete tuples/rows

- `c.execute("“ DELETE FROM Student  
WHERE name = “Harper Taylor”  
”“)`

	0	1	2
0	201921323	Ava Smith	2
1	201832220	Ben Johnson	3
2	202003219	Charlie Jones	1
3	202045234	Dan Norris	1
4	201985603	Emily Wood	1
5	201933222	Freddie Harris	2
6	201875940	Grace Clarke	2

- Disconnecting from database
  - `conn.close()`

## ***Querying Databases in Python***

- Get query results in data frame
  - ```
q1 = c.execute("SELECT final_mark
                  FROM Grade
                  WHERE course_id = 'ST101'
                  ").fetchall()
```
 - `pd.DataFrame(q1)`

| | 0 |
|---|----|
| 0 | 78 |
| 1 | 60 |
| 2 | 47 |

- Get results in alphabetical order
 - ```
q2 = c.execute("SELECT Student.name
 FROM Grade, Student
 WHERE Grade.course_id='ST101' AND
 Student.student_id=Grade.student_id
 ").fetchall()
```

ORDER BY Student.name

'''").fetchall()

o pd.DataFrame(q2)

```
0
0 Ava Smith
1 Charlie Jones
2 Emily Wood
```

- Get distinct results

o q3 = c.execute("““ SELECT DISTINCT Course.name  
FROM Student, Grade, Course  
WHERE (Student.name = ‘Ava Smith’ OR  
Student.name = ‘Freddie Harris’) AND  
Student.student\_id = Grade.student\_id  
AND Course.course\_id = Grade.course\_id  
”““).fetchall()

o pd.DataFrame(q3)

```
0
0 programming for data science
1 Managing and Visualising Data
2 Databases
```

- Get calculated results

o q4 = c.execute("““ SELECT course\_id, AVG(final\_mark)  
AS avg\_mark  
FROM Grade  
GROUP BY course\_id  
”““).fetchall()

o pd.DataFrame(q4)

	0	1
0	ST101	61.666667
1	ST115	82.333333
2	ST207	66.500000

## Useful Resources

- - <http://>